

Assignment is at the bottom!

```
In [1]: from sklearn.linear_model import LogisticRegression
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import numpy as np


from pylab import rcParams
rcParams['figure.figsize'] = 20, 10

from sklearn.linear_model import LogisticRegression as Model
```

```
In [2]: y = np.concatenate([np.zeros(10), np.ones(10)])
x = np.linspace(0, 10, len(y))
```

```
In [3]: plt.scatter(x, y, c=y)
```

```
Out[3]: <matplotlib.collections.PathCollection at 0x144d5a0a8f0>
```




```
In [4]: model = LogisticRegression()
```

```
In [5]: model.fit(x.reshape(-1, 1), y)
```

```
Out[5]: ▾ LogisticRegression
LogisticRegression()
```

```
In [6]: plt.scatter(x, y, c=y)
plt.plot(x, model.predict_proba(x.reshape(-1, 1))[:, 1])
```

```
Out[6]: [<matplotlib.lines.Line2D at 0x144d83303a0>]
```




```
In [7]: b, b0 = model.coef_, model.intercept_
        model.coef_, model.intercept_
```

```
Out[7]: (array([[1.46709085]]), array([-7.33542562]))
```

```
In [8]: plt.plot(x, 1/(1+np.exp(-x)))
```

```
Out[8]: [<matplotlib.lines.Line2D at 0x144d88a56c0>]
```




```
In [9]: b
```

```
Out[9]: array([[1.46709085]])
```

```
In [10]: plt.plot(x, 1/(1+np.exp(-(b[0]*x + b0))))
```

```
Out[10]: [<matplotlib.lines.Line2D at 0x144d8922b90>]
```




```
In [11]: from mpl_toolkits.mplot3d import Axes3D # noqa: F401 unused import

import matplotlib.pyplot as plt
from matplotlib import cm
from matplotlib.ticker import LinearLocator, FormatStrFormatter
import numpy as np

fig = plt.figure()
ax = fig.add_subplot(projection='3d') # gca wasn't working

# Make data.
X = np.arange(-10, 10, 0.25)
Y = np.arange(-10, 10, 0.25)
X, Y = np.meshgrid(X, Y)
R = np.sqrt(X**2 + Y**2)
Z = 1/(1+np.exp(-(b[0]*X + b[0]*Y + b0)))
surf = ax.plot_surface(X, Y, Z, cmap=cm.coolwarm,
                      linewidth=0, antialiased=False)
```



In [12]: X

```
Out[12]: array([[ -10.   ,  -9.75,  -9.5 , ...,   9.25,   9.5 ,   9.75],
                [ -10.   ,  -9.75,  -9.5 , ...,   9.25,   9.5 ,   9.75],
                [ -10.   ,  -9.75,  -9.5 , ...,   9.25,   9.5 ,   9.75],
                ...,
                [ -10.   ,  -9.75,  -9.5 , ...,   9.25,   9.5 ,   9.75],
                [ -10.   ,  -9.75,  -9.5 , ...,   9.25,   9.5 ,   9.75],
                [ -10.   ,  -9.75,  -9.5 , ...,   9.25,   9.5 ,   9.75]])
```

In [13]: Y


```
Out[13]: array([[ -10.   ,  -10.   ,  -10.   , ...,  -10.   ,  -10.   ,  -10.   ],
                [  -9.75,  -9.75,  -9.75, ...,  -9.75,  -9.75,  -9.75],
                [  -9.5 ,  -9.5 ,  -9.5 , ...,  -9.5 ,  -9.5 ,  -9.5 ],
                ...,
                [   9.25,   9.25,   9.25, ...,   9.25,   9.25,   9.25],
                [   9.5 ,   9.5 ,   9.5 , ...,   9.5 ,   9.5 ,   9.5 ],
                [   9.75,   9.75,   9.75, ...,   9.75,   9.75,   9.75]])
```

What if the data doesn't really fit this pattern?

```
In [14]: y = np.concatenate([np.zeros(10), np.ones(10), np.zeros(10)])
x = np.linspace(0, 10, len(y))
```

```
In [15]: plt.scatter(x,y, c=y)
```

```
Out[15]: <matplotlib.collections.PathCollection at 0x144d8bbdc90>
```




```
In [16]: model.fit(x.reshape(-1, 1),y)
```

```
Out[16]: ▾ LogisticRegression
LogisticRegression()
```

```
In [17]: plt.scatter(x,y)
plt.plot(x, model.predict_proba(x.reshape(-1, 1)))
```

```
Out[17]: [<matplotlib.lines.Line2D at 0x144d8ce3760>,
<matplotlib.lines.Line2D at 0x144d8ce37c0>]
```



```
In [18]: model1 = LogisticRegression()
model1.fit(x[:15].reshape(-1, 1), y[:15])
```


```
Out[18]: ▾ LogisticRegression
LogisticRegression()
```

```
In [19]: model2 = LogisticRegression()
model2.fit(x[15:].reshape(-1, 1), y[15:])
```

```
Out[19]: ▾ LogisticRegression
LogisticRegression()
```

```
In [20]: plt.scatter(x, y, c=y)
plt.plot(x, model1.predict_proba(x.reshape(-1, 1))[:, 1] * model2.predict_proba(x.reshape(-1, 1))[:, 1])
```

```
Out[20]: [<matplotlib.lines.Line2D at 0x144d5a9de70>]
```



```
In [21]: df = pd.read_csv('../data/adult.data', index_col=False)
golden = pd.read_csv('../data/adult.test', index_col=False)
```

```
In [22]: from sklearn import preprocessing

enc = preprocessing.OrdinalEncoder()
```

```
In [23]: transform_columns = ['sex', 'workclass', 'education', 'marital-status',
                             'occupation', 'relationship', 'race', 'sex',
                             'native-country', 'salary']
```

```
In [24]: x = df.copy()

x[transform_columns] = enc.fit_transform(df[transform_columns])

golden['salary'] = golden.salary.replace(' <=50K.', ' <=50K').replace(' >50K.', ' >50K')
xt = golden.copy()

xt[transform_columns] = enc.transform(golden[transform_columns])
```

```
In [25]: df.salary.unique()
```

```
Out[25]: array([' <=50K', ' >50K'], dtype=object)
```

```
In [26]: golden.salary.replace(' <=50K.', ' <=50K').replace(' >50K.', ' >50K').unique()
```

```
Out[26]: array([' <=50K', ' >50K'], dtype=object)
```

```
In [27]: model.fit(preprocessing.scale(x.drop('salary', axis=1)), x.salary)
```

```
Out[27]: LogisticRegression
LogisticRegression()
```

```
In [28]: pred = model.predict(preprocessing.scale(x.drop('salary', axis=1)))
pred_test = model.predict(preprocessing.scale(xt.drop('salary', axis=1)))
```

In [29]: `x.head()`

Out[29]:

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex	capital-gain
0	39	7.0	77516	9.0	13	4.0	1.0	1.0	4.0	1.0	2.0
1	50	6.0	83311	9.0	13	2.0	4.0	0.0	4.0	1.0	0.0
2	38	4.0	215646	11.0	9	0.0	6.0	1.0	4.0	1.0	0.0
3	53	4.0	234721	1.0	7	2.0	6.0	0.0	2.0	1.0	0.0
4	28	4.0	338409	9.0	13	2.0	10.0	5.0	2.0	0.0	0.0

In [30]: `from sklearn.metrics import (
 accuracy_score,
 classification_report,
 confusion_matrix, auc, roc_curve
)`

In [31]: `accuracy_score(x.salary, pred)`

Out[31]: 0.8250360861152913

In [32]: `confusion_matrix(x.salary, pred)`

Out[32]: `array([[23300, 1420],
 [4277, 3564]], dtype=int64)`

In [33]: `print(classification_report(x.salary, pred))`

```

              precision    recall  f1-score   support

    0.0         0.84        0.94        0.89        24720
    1.0         0.72        0.45        0.56         7841

 accuracy                   0.83        32561
 macro avg              0.78        0.70        0.72        32561
 weighted avg           0.81        0.83        0.81        32561

```

In [34]: `print(classification_report(xt.salary, pred_test))`

```

              precision    recall  f1-score   support

    0.0         0.85        0.94        0.89        12435
    1.0         0.70        0.45        0.55         3846

 accuracy                   0.82        16281
 macro avg              0.77        0.69        0.72        16281
 weighted avg           0.81        0.82        0.81        16281

```

Assignment

1. Use your own dataset (Heart.csv is acceptable), create a train and a test set, and build 2 models: Logistic Regression and Decision Tree (shallow). Compare the test results using classification_report and confusion_matrix. Explain which algorithm is optimal

2. Repeat 1. but let the Decision Tree be much deeper to allow over-fitting. Compare the two models' test results again, and explain which is optimal

```
In [35]: from sklearn.svm import LinearSVC
from sklearn.linear_model import LogisticRegression
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
sns.set(font_scale=1.5)
import numpy as np

from pylab import rcParams
rcParams['figure.figsize'] = 20, 10

from sklearn.linear_model import LogisticRegression as Model
```

```
In [36]: # Logistic Regression
heart = pd.read_csv('../data/Heart.csv') # Had to remove extra '/../'
heart = pd.DataFrame(heart)
heart.dropna(inplace=True)
heart['AHD'] = heart['AHD'].map({'Yes': 1, 'No': 0})
# feature_cols = ['RestBP', 'Chol']
# feature_columns = ['RestBP', 'Chol']
x = heart.drop(['AHD', 'ChestPain', 'Thal'], axis=1)
y = heart.AHD
heart['AHD'] = heart['AHD'].map({'Yes': 1, 'No': 0})
```

```
In [37]: # Train & Test Sets
from sklearn.model_selection import train_test_split
xtrain, xtest, ytrain, ytest = train_test_split(x, y, test_size=0.2)
(
    xtrain.shape,
    # xtest.shape,
    ytrain.shape #,
    # ytest.shape
)
```

```
Out[37]: ((237, 12), (237,))
```

```
In [38]: # from sklearn.preprocessing import StandardScaler # Trouble with Reshape, using a sin
# normalizer = StandardScaler()
# [['RestBP', 'Chol']] = normalizer.fit_transform(xtrain)
ytrain
```

```
Out[38]: 56      1
237      1
111      1
162      0
34       0

..
79       1
104      1
32       1
253      0
78       0
Name: AHD, Length: 237, dtype: int64
```

```
In [39]: model = LogisticRegression()
model.fit(xtrain,ytrain)
lr_predictions = model.predict(xtest)
lr_predictions
pd.array(ytest)
```

C:\Users\Brett\anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

n_iter_i = _check_optimize_result(

```
Out[39]: <PandasArray>
[0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0, 1, 0, 1,
 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 1, 0, 0, 1, 0, 1, 0, 1, 0,
 0, 0, 1, 1, 0, 1, 1, 1]
Length: 60, dtype: int64
```

```
In [40]: # Classification Report
print(classification_report(pd.array(ytest), lr_predictions))
# Confusion Matrix
print(confusion_matrix(pd.array(ytest), lr_predictions))
```

	precision	recall	f1-score	support
0	0.80	0.76	0.78	37
1	0.64	0.70	0.67	23
accuracy			0.73	60
macro avg	0.72	0.73	0.72	60
weighted avg	0.74	0.73	0.74	60

```
[[28  9]
 [ 7 16]]
```

```
In [41]: # Decision Tree
from sklearn.tree import DecisionTreeClassifier
decisiontreemodel = DecisionTreeClassifier(criterion='entropy', max_depth= 3)
decisiontreemodel.fit(xtrain, ytrain)
```

```
dt_predictions = decisiontreemodel.predict(xtest)
dt_predictions
```

```
Out[41]: array([0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 1, 1,
        0, 1, 0, 1, 0, 1, 0, 1, 1, 0, 1, 0, 0, 1, 0, 1, 1, 0, 1, 1, 1, 0,
        0, 0, 1, 1, 1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 1, 1], dtype=int64)
```

```
In [42]: # Classification Report
print(classification_report(pd.array(ytest), dt_predictions))
# Confusion Matrix
print(confusion_matrix(pd.array(ytest), dt_predictions))
```

	precision	recall	f1-score	support
0	0.79	0.70	0.74	37
1	0.59	0.70	0.64	23
accuracy			0.70	60
macro avg	0.69	0.70	0.69	60
weighted avg	0.71	0.70	0.70	60

```
[[26 11]
 [ 7 16]]
```

```
In [43]: # Comparison
# As seen in the classification reports, the logistic regression is superior in precis
```

```
In [44]: # Decision Tree - Deeper
decisiontreemodel_deep = DecisionTreeClassifier(criterion='entropy', max_depth= 9999)
decisiontreemodel_deep.fit(xtrain, ytrain)
dt_predictions_deep = decisiontreemodel_deep.predict(xtest)
dt_predictions_deep
```

```
Out[44]: array([1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 1, 0, 0, 1, 1, 1,
        0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 1, 1, 0, 1, 0, 1, 0, 0, 1, 1, 0, 0,
        0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 1, 1, 0, 0, 1], dtype=int64)
```

```
In [45]: # Classification Report
print(classification_report(pd.array(ytest), dt_predictions_deep))
# Confusion Matrix
print(confusion_matrix(pd.array(ytest), dt_predictions_deep))
```

	precision	recall	f1-score	support
0	0.72	0.70	0.71	37
1	0.54	0.57	0.55	23
accuracy			0.65	60
macro avg	0.63	0.63	0.63	60
weighted avg	0.65	0.65	0.65	60

```
[[26 11]
 [10 13]]
```

```
In [ ]: # By increasing the depth to a point where the decision tree overfits, we see decrease
# The model is too closely trained on the training set, and not adaptable to other dat
```

```
In [46]: (
    decisiontreemodel.tree_.node_count,
```

```
decisiontreemodel_deep.tree_.node_count  
)
```

Out[46]: (15, 71)