

Neural Networks - Image Recognition

```
In [1]: import keras
        from keras.datasets import mnist
        from keras.models import Sequential
        from keras.optimizers import RMSprop
        from keras.layers import Dense, Dropout, Flatten
        from keras.layers import Conv2D, MaxPooling2D
        from keras import backend
```

```
In [2]: import matplotlib.pyplot as plt
        %matplotlib inline
```

```
In [3]: # the data, shuffled and split between train and test sets
        (x_train, y_train), (x_test, y_test) = mnist.load_data()

        x_train = x_train.reshape(60000, 784)
        x_test = x_test.reshape(10000, 784)
        x_train = x_train.astype('float32')
        x_test = x_test.astype('float32')
        x_train /= 255
        x_test /= 255
        print(x_train.shape[0], 'train samples')
        print(x_test.shape[0], 'test samples')
```

60000 train samples
10000 test samples

1. Add random noise (see below on `size` parameter on `np.random.normal`) to the images in training and testing. **Make sure each image gets a different noise feature added to it. Inspect by printing out several images. Note - the `size` parameter should match the data.**

```
In [4]: # Noise is added here

        # Generate random noise with the same shape as the image
        import numpy as np
        mean = .5
        stddev = .16
        noise = np.random.normal(mean, stddev, x_train[1].shape)

        # Add the noise to the image
        image_noise_added = x_train[1] + noise
        image_noise_added


        # The max value of the noise should not grossly surpass 1.0
        max(noise)
```

Out[4]: 0.9814309422337107

```
In [5]: # Max of noise is less than .1002; very slightly over 1
```


```
In [6]: # New Image vs Old  
# New Image  
plt.imshow(image_noise_added.reshape(28, 28))
```

```
Out[6]: <matplotlib.image.AxesImage at 0x27146484400>
```



```
In [7]: # Old Image  
plt.imshow(x_train[1].reshape(28, 28))
```

```
Out[7]: <matplotlib.image.AxesImage at 0x271484f4760>
```



```
In [8]: ### Add noise to whole set of images
# Noise is added here

# Generate random noise with the same shape as the image
import numpy as np
mean = .5
stddev = .16
noise = np.random.normal(mean, stddev, x_train.shape)

# Generate random noise with the same shape as the image
test_noise = np.random.normal(mean, stddev, x_test.shape)
y_train_noise = np.random.normal(mean, stddev, y_train.shape)
y_test_noise = np.random.normal(mean, stddev, y_test.shape)

# Add noise
x_train_noise = x_train + noise
x_test_noise = x_test + test_noise
y_train_noise = y_train + y_train_noise
y_test_noise = y_test + y_test_noise

# Add noise
x_test_noise[1]
```

```

Out[8]: array([ 0.57687232,  0.40826787,  0.48698683,  0.50217036,  0.645252  ,
 0.61910315,  0.41440629,  0.51303236,  0.63962366,  0.64682592,
 0.37467939,  0.52838516,  0.57147313,  0.60116104,  0.42979264,
 0.47278035,  0.17707505,  0.6037902 ,  0.38793955,  0.4623349 ,
 0.19363086,  0.37676522,  0.67193465,  0.58154555,  0.46630765,
 0.34852071,  0.42603748,  0.60274664,  0.34968219,  0.57033017,
 0.32556246,  0.34093753,  0.46699395,  0.40841534,  0.31135967,
 0.40805849,  0.55859174,  0.638774  ,  0.30170078,  0.73114616,
 0.37730497,  0.63878796,  0.44348356,  0.5813878 ,  0.60068954,
 0.28910888,  0.62700597,  0.36169634,  0.46364495,  0.46958263,
 0.50642522,  0.4631466 ,  0.48307732,  0.63050008,  0.54317324,
 0.46467005,  0.76507951,  0.41156469,  0.67368421,  0.45407504,
 0.41759021,  0.26744614,  0.52759676,  0.53368195,  0.48003363,
 0.51910648,  0.79028283,  0.72157947,  0.47613033,  0.38042947,
 0.71802766,  0.35599436,  0.50874887,  0.33294799,  0.33277805,
 0.51121625,  0.29037356,  0.54504028,  0.80172431,  0.65532005,
 0.66518034,  0.47198225,  0.36786903,  0.6233744 ,  0.42654806,
 0.626913  ,  0.39782186,  0.42302727,  0.50334411,  0.4370632 ,
 0.4251654 ,  0.30500427,  0.55564946,  0.38227468,  0.96422872,
 0.91836584,  1.24052622,  1.46307529,  1.56094511,  1.25740408,
 0.69438408,  0.6005917 ,  0.53072081,  0.45989333,  0.31914228,
 0.47099433,  0.22500937,  0.5613782 ,  0.75790809,  0.70932527,
 0.45917871,  0.39644749,  0.33798038,  0.32350235,  0.61883753,
 0.39654337,  0.70952601,  0.60944194,  0.50570285,  0.56547912,
 0.7316817 ,  1.48010999,  1.20563341,  1.5354358 ,  1.30711049,
 1.5711794 ,  1.43229737,  1.29984594,  1.38860626,  0.44369197,
 0.34746496,  0.54786675,  0.46418887,  0.72041538,  0.71813562,
 0.34223791,  0.45970136,  1.14808828,  0.47099867,  0.27251569,
 0.55721335,  0.38836619,  0.29644195,  0.41930418,  0.59695966,
 0.53640419,  0.64318424,  0.35857606,  1.12010692,  1.56477894,
 1.58059331,  1.60154089,  1.27910914,  1.04854077,  1.45419382,
 1.60501054,  1.61817443,  0.91370378,  0.5173893 ,  0.58896817,
 0.39474733,  0.58212097,  0.46727581,  0.79173531,  0.60242163,
 0.45962592,  0.54456266,  0.37801249,  0.56099927,  0.56874342,
 0.6159749 ,  0.50175787,  0.55086682,  0.6257764 ,  0.75945265,
 1.07127152,  1.36776284,  1.35008523,  1.36692701,  0.77233484,
 0.57868446,  0.30944101,  0.71401451,  1.20846498,  1.65998397,
 1.29847969,  0.66515439,  0.36115685,  0.22984019,  0.63827504,
 0.51871401,  0.41042776,  0.57781973,  0.4880221 ,  0.77902976,
 0.67478363,  0.57619222,  0.29292569,  0.3714438 ,  0.47110892,
 0.65160108,  0.67104572,  0.39016473,  0.91870886,  1.25853786,
 1.43544991,  0.73563021,  0.78235505,  0.38323818,  0.58749505,
 0.78000725,  1.69187491,  1.49744226,  0.88697527,  0.79193357,
 0.56011929,  0.63032284,  0.69932626,  0.37608275,  0.61027271,
 0.52195608,  0.48601603,  0.74008066,  0.40982626,  0.61027295,
 0.47972026,  0.3698776 ,  0.55671247,  0.61515198,  0.40591306,
 0.44684098,  0.5449375 ,  0.51541301,  0.4153461 ,  0.46303803,
 0.50765033,  0.57844954,  0.42400258,  1.22140617,  1.61883959,
 1.42202093,  0.58517293,  0.23959164,  0.52297909,  0.64580432,
 0.34583476,  0.48682419,  0.30327991,  1.04286853,  0.31668652,
 0.59552655,  0.43303125,  0.60815892,  0.21081781,  0.64153105,
 0.64774413,  0.71064805,  0.42262043,  0.40247365,  0.44596662,
 0.40861889,  0.49788938,  0.52451359,  0.62464182,  0.39184073,
 1.34829301,  1.70225717,  1.53762747,  1.51560417,  0.48299895,
 0.33936687,  0.69238985,  0.32186298,  0.2586233 ,  0.28474299,
 0.55044112,  0.52029234,  0.46978064,  0.66648489,  0.32801189,
 0.41059155,  0.58940078,  0.67963343,  0.52471417,  0.61162958,
 0.80035667,  0.66476024,  0.458238 ,  0.75578721,  0.56674768,
 0.3457525 ,  0.72454083,  0.93367028,  1.49255573,  1.35476919,
 1.32076445,  1.14075649,  0.40289053,  0.38107664,  0.444363  ,

```

0.57212264,	0.38351332,	0.68051138,	0.54716453,	0.62290098,
0.25550828,	0.70895076,	0.64530895,	0.27380722,	0.77455907,
0.50103175,	0.58251633,	0.41462832,	0.14251818,	0.46168851,
0.60593235,	0.66388161,	0.46078412,	0.57685837,	0.30246699,
1.2257414 ,	1.3880618 ,	1.43614418,	0.95321121,	0.45152506,
0.19778098,	0.54767181,	0.29674595,	0.41591414,	0.52391182,
0.57302419,	0.58286705,	0.676693 ,	0.45983768,	0.70680645,
0.47206472,	0.68424349,	0.13648687,	0.5869302 ,	0.55534703,
0.48605483,	0.34212007,	0.56896385,	0.60127747,	0.36992377,
0.44773107,	0.69464795,	1.10740255,	1.43281916,	1.40614555,
1.19718594,	0.56582603,	0.21850394,	0.5248323 ,	0.5604059 ,
0.49926849,	0.49443567,	0.46567549,	0.42704746,	0.60985869,
0.61681303,	0.7036364 ,	0.29623638,	0.50978999,	0.63400692,
0.41059443,	0.30164586,	0.55097491,	0.69367964,	0.44303171,
0.58091256,	0.44285741,	0.56410263,	0.61156829,	0.75049957,
1.43629381,	1.52172267,	1.34290132,	0.55388268,	0.38844558,
0.48508947,	0.3696634 ,	0.63216608,	0.48751598,	0.37656868,
0.43954913,	0.23700829,	0.4159265 ,	0.41542122,	0.47840707,
0.44554244,	0.72638413,	0.56213323,	0.76256056,	0.44031373,
0.43894338,	0.25594314,	0.33465653,	0.44415611,	0.51582468,
0.55120162,	0.66069936,	1.31538128,	1.66328685,	1.44516188,
1.02270407,	0.12003318,	0.69552721,	0.53139921,	0.33769814,
0.79705492,	0.75478203,	0.34682688,	0.70877293,	0.79783351,
0.36799432,	0.72164162,	0.46368178,	0.61005969,	0.76968896,
0.80616255,	0.46435952,	0.6975923 ,	0.59662083,	0.6061286 ,
0.33113977,	0.46146188,	0.62182404,	0.4812855 ,	0.75222937,
1.2436928 ,	1.78719514,	1.31072272,	0.3970764 ,	0.31548864,
0.56647389,	0.49042147,	0.49246477,	0.70088437,	0.32063056,
0.43199672,	0.4218937 ,	0.60392227,	0.65387342,	0.44636835,
0.61034481,	0.32493709,	0.44398562,	0.63548597,	0.47206533,
0.58278564,	0.52560018,	0.46213753,	0.53412207,	0.61547489,
0.5045515 ,	0.61784269,	1.20916063,	1.31296184,	1.53725133,
1.17186029,	0.34904539,	0.50627498,	0.50687739,	0.53640772,
0.56586986,	0.43075534,	0.60988961,	0.62925669,	0.56843169,
0.61100649,	0.43874335,	0.53837976,	0.62260578,	0.41965348,
0.6808537 ,	0.24438297,	0.49623488,	0.57181072,	0.34635712,
0.50992352,	0.5679127 ,	0.47805241,	0.50743278,	1.05503901,
1.51676278,	1.47122376,	1.43540835,	0.68739887,	0.81629148,
0.3794722 ,	0.55767566,	0.41346537,	0.49601173,	0.47879598,
0.66224115,	0.41859635,	0.63192803,	0.61052622,	0.37330411,
0.47443541,	0.28885453,	0.56104045,	0.48758825,	0.68472919,
0.7342782 ,	-0.0146911 ,	0.74080096,	0.44821362,	0.80058136,
0.56618772,	0.48387254,	1.40739311,	1.36079843,	1.39607238,
0.50969077,	0.5432472 ,	0.49923699,	0.541422 ,	0.3493232 ,
0.53854721,	0.52690456,	0.36931214,	0.60147836,	0.56491768,
0.6680341 ,	0.58817732,	0.49161062,	0.5382499 ,	0.66288018,
0.24496488,	0.60784583,	0.5100051 ,	0.45461803,	0.27164354,
0.40266897,	0.50877334,	0.63210122,	0.57207429,	0.51167262,
1.58497908,	1.16691554,	1.36532609,	1.02976064,	0.68604067,
0.6000592 ,	0.64854475,	0.50860073,	0.56702268,	0.39755904,
0.50356904,	0.52461119,	0.61202455,	0.75421722,	0.81392503,
1.2825361 ,	1.48882145,	1.38675446,	0.33056584,	0.21335745,
0.14683961,	0.64504509,	0.63019434,	0.56613274,	0.74820247,
0.71333899,	0.61619546,	0.07963843,	1.31058875,	1.61219212,
1.60755701,	1.45311576,	1.47463772,	1.4965448 ,	1.58811844,
1.74563024,	1.02335638,	0.93359924,	0.88620072,	1.33135528,
1.26372337,	1.6930712 ,	1.56824574,	1.5644695 ,	1.28664689,
1.39342324,	0.97731335,	0.59283548,	0.47171893,	0.45035895,
0.65857503,	0.52476677,	0.51760645,	0.6293517 ,	0.59639758,
0.57505975,	1.25416421,	1.70952544,	1.39940906,	1.32892249,

```

1.46538103, 1.55307087, 1.59780653, 1.66271572, 1.24499921,
1.53049975, 1.30418381, 1.54324893, 1.59013243, 1.25245885,
1.48635307, 0.99729451, 1.08297132, 1.12765557, 0.68652973,
0.77951035, 0.44624082, 0.40428207, 0.64056763, 0.72175819,
0.38757577, 0.2435911 , 0.69577388, 0.53197946, 0.50748117,
0.92806069, 1.03204323, 0.78733256, 1.20431391, 0.96670428,
1.76335259, 1.56037524, 1.45832627, 1.19755596, 0.90655534,
1.00699761, 0.40481721, 0.31130171, 0.57530393, 0.38067106,
0.2374199 , 0.30717348, 0.59960756, 0.50586556, 0.42062948,
0.48058779, 0.32149952, 0.47543075, 0.17546164, 0.77975522,
0.5097491 , 0.67831857, 0.55793433, 0.36801297, 0.19546406,
0.24104791, 0.74321123, 0.30950039, 0.36157467, 0.6576656 ,
0.33552428, 0.36398714, 0.44787139, 0.39748342, 0.45472615,
0.69077358, 0.44023924, 0.49691089, 0.55107735, 0.60205863,
0.51575186, 0.40972205, 0.5528749 , 0.7018578 , 0.29692798,
0.56124433, 0.74747962, 0.77443373, 0.48150657, 0.52666454,
0.97278558, 0.34108848, 0.36858934, 0.30622875, 0.55314357,
0.4364178 , 0.65828641, 0.47318976, 0.87802943, 0.50157545,
0.58784343, 0.14233734, 0.63692678, 0.27199073, 0.33432951,
0.55719796, 0.01497966, 0.51575917, 0.22319883, 0.18286803,
0.59889024, 0.53215933, 0.46862034, 0.7253405 , 0.51431413,
0.51807683, 0.352295 , 0.72957085, 0.24179591, 0.3677939 ,
0.59017383, 0.4299084 , 0.57120308, 0.54409436, 0.13898291,
0.30659389, 0.32858998, 0.75806143, 0.25855974, 0.51663139,
0.63212019, 0.40535639, 0.33615966, 0.74985285, 0.41696861,
0.42385191, 0.27142034, 0.59889696, 0.53467767, 0.58737025,
0.8085762 , 0.43750174, 0.78269892, 0.61612576, 0.78229204,
0.20155359, 0.49406308, 0.69594044, 0.90704146, 0.48068657,
0.2593931 , 0.49739493, 0.55527354, 0.2703061 , 0.52558612,
0.65110954, 0.84393377, 0.41588745, 0.68644318, 0.47088899,
0.63850876, 0.61065792, 0.53675139, 0.99290694, 0.35616339,
0.5933731 , 0.32941645, 0.43484905, 0.63120649, 0.62207694,
0.45278085, 0.26114508, 0.39601245, 0.54658679, 0.45476478,
0.60711821, 0.55012664, 0.42348607, 0.76234562, 0.44878522,
0.60790947, 0.3595148 , 0.23428315, 0.54523711, 0.63661374,
0.05622986, 0.49892666, 0.32963783, 0.52101281, 0.59030581,
0.3802501 , 0.55527802, 0.53023526, 0.73779477])

```

```


In [9]: # Compare Noises to Ensure they are Different
# Noise on Image 122
plt.imshow(noise[122].reshape(28, 28))

```

```


Out[9]: <matplotlib.image.AxesImage at 0x27148566200>

```



```
In [10]: # Noise on Image 124  
plt.imshow(noise[124].reshape(28, 28))
```

```
Out[10]: <matplotlib.image.AxesImage at 0x271485de350>
```



```
In [11]: # As the two visual representations (images) of the noise show different patterns, the
# for each image
```

1. Compare the **accuracy** of train and val after N epochs for MLNN with and without noise.

```
In [12]: %%capture
## Code From Base for Creating Neural Network Without Noise
batch_size = 128
num_classes = 10
epochs = 20

# convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)

model = Sequential()
model.add(Dense(512, activation='relu', input_shape=(784,)))
model.add(Dropout(0.2))
model.add(Dense(512, activation='relu'))
model.add(Dropout(0.2))
model.add(Dense(10, activation='softmax'))

model.summary()

model.compile(loss='categorical_crossentropy',
              optimizer="adam",
              metrics=['accuracy'])

history = model.fit(x_train, y_train,
                   batch_size=batch_size,
                   epochs=epochs,
                   verbose=1,
                   validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

```
In [13]: import pandas as pd
(
np.max(x_train_noise),
np.min(x_train_noise),
np.max(x_train),
np.min(x_train)
)
```

```
Out[13]: (2.317632454193437, -0.4185654653815156, 1.0, 0.0)
```

```
In [14]: # Normalization
mean = np.mean(x_train_noise)
std = np.std(x_train_noise)
x_train_noise_norm = (x_train_noise - mean) / std

mean = np.mean(x_test_noise)
std = np.std(x_test_noise)
x_test_noise_norm = (x_test_noise - mean) / std

mean = np.mean(y_train_noise)
```



```
std = np.std(y_train_noise)
y_train_noise_norm = (y_train_noise - mean) / std

mean = np.mean(y_test_noise)
std = np.std(y_test_noise)
y_test_noise_norm = (y_test_noise - mean) / std
```

```
In [15]: %%capture
        ## Updated Code for Creating Neural Network With Noise
        batch_size = 128
        num_classes = 10
        # epochs = 20 ## Commenting out; this ensures the same number of epochs is used for bc

        # convert class vectors to binary class matrices - using noise added images
        y_train_bcm = keras.utils.to_categorical(y_train_noise_norm, num_classes)
        y_test_bcm = keras.utils.to_categorical(y_test_noise_norm, num_classes)

        noisy_model = Sequential()
        noisy_model.add(Dense(512, activation='relu', input_shape=(784,)))
        noisy_model.add(Dropout(0.2))
        noisy_model.add(Dense(512, activation='relu'))
        noisy_model.add(Dropout(0.2))
        noisy_model.add(Dense(10, activation='softmax'))

        noisy_model.summary()

        noisy_model.compile(loss='categorical_crossentropy',
                            optimizer="adam",
                            metrics=['accuracy'])

        history = noisy_model.fit(x_train_noise_norm, y_train_bcm,
                                   batch_size=batch_size,
                                   epochs=epochs,
                                   verbose=1,
                                   validation_data=(x_test_noise_norm, y_test_bcm))
        noisy_score = noisy_model.evaluate(x_test_noise_norm, y_test_bcm, verbose=0)
        print('Test loss:', noisy_score[0])
        print('Test accuracy:', noisy_score[1])
```

```
In [16]: print('Test loss:', score[0]) # Test loss & accuracy statistics for the model trained
        print('Test accuracy:', score[1])
        print('Test loss:', noisy_score[0]) # The same statistics for the model trained with r
        print('Test accuracy:', noisy_score[1])
```

```
Test loss: 0.0920642837882042
Test accuracy: 0.9800999760627747
Test loss: 0.1428346335887909
Test accuracy: 0.9804999828338623
```

```
In [17]: # Comparing the without-noise and with-noise models, we find that the noise reduces ac
        # a training series lasting n=20 epochs. Additionally, the test loss statistic is noic
```

1. Vary the amount of noise by changing the `scale` parameter in `np.random.normal` by a factor. Use `.1`, `.5`, `1.0`, `2.0`, `4.0` for the `scale` and keep track of the `accuracy` for training and validation and plot these results.

```

In [18]: %%capture
## Combining Noise Code with Model-Running code for a single, repeatable block:
# Listing Scale (i.e. Standard Deviation) Values over which to train models
stddev_values = [0.1, 0.5, 1.0, 2.0, 4.0] # Values given
# Initializing Lists
loss_values = []
accuracy_values = []

# Loop:
for stddev in stddev_values:

    # Generate random noise with the same shape as the image
    mean = .5
    stddev = .16

    # Generate random noise with the same shape as the image
    x_train_noise = np.random.normal(mean, stddev, x_train.shape)
    x_test_noise = np.random.normal(mean, stddev, x_test.shape)
    y_train_noise = np.random.normal(mean, stddev, y_train.shape)
    y_test_noise = np.random.normal(mean, stddev, y_test.shape)

    # Add noise
    x_train_noise = x_train + x_train_noise
    x_test_noise = x_test + x_test_noise
    y_train_noise = y_train + y_train_noise
    y_test_noise = y_test + y_test_noise

    ## Updated Code for Creating Neural Network With Noise
    batch_size = 128
    num_classes = 10
    epochs = 20 # Always using same number of epochs

    # convert class vectors to binary class matrices
    y_train_bcm = keras.utils.to_categorical(y_train_noise_norm, num_classes)
    y_test_bcm = keras.utils.to_categorical(y_test_noise_norm, num_classes)

    noisy_model = Sequential()
    noisy_model.add(Dense(512, activation='relu', input_shape=(784,)))
    noisy_model.add(Dropout(0.2))
    noisy_model.add(Dense(512, activation='relu'))
    noisy_model.add(Dropout(0.2))
    noisy_model.add(Dense(10, activation='softmax'))

    noisy_model.summary()

    noisy_model.compile(loss='categorical_crossentropy',
                        optimizer="adam",
                        metrics=['accuracy'])

    history = noisy_model.fit(x_train_noise_norm, y_train_bcm,
                              batch_size=batch_size,
                              epochs=epochs,
                              verbose=1,
                              validation_data=(x_test_noise_norm, y_test_bcm))

    noisy_score = noisy_model.evaluate(x_test_noise_norm, y_test_bcm, verbose=0)
    print('Test loss:', noisy_score[0])
    print('Test accuracy:', noisy_score[1])
    # Append Loss & Accuracy To List
    test_loss = noisy_score[0]

```

```
test_accuracy = noisy_score[1]
loss_values.append(test_loss)
accuracy_values.append(test_accuracy)
```

```
In [19]: # Compile Loss and Accuracy statistics for the models into a single data frame
accuracy_loss_results_dataframe = pd.DataFrame({'stddev_values': stddev_values, 'loss_
accuracy_loss_results_dataframe
```

```
Out[19]:
```

	stddev_values	loss_values	accuracy_values
0	0.1	0.161019	0.9780
1	0.5	0.135908	0.9786
2	1.0	0.144020	0.9772
3	2.0	0.136160	0.9782
4	4.0	0.150436	0.9775

```
In [20]: # Plotting Results
fig, ax1 = plt.subplots()
loss_line = ax1.plot(stddev_values, loss_values, color = 'green', label = 'Test Loss S

# Loss Line Plot
ax1.set_xlabel('Scale (Standard Deviation)')
ax1.set_ylabel('Loss Values')
ax1.set_title('Combined Loss and Accuracy vs Scale Plot')

# Accuracy Line Plot
ax2 = ax1.twinx()
accuracy_line = ax2.plot(stddev_values, accuracy_values, color='orange', label='Test A
ax2.set_ylabel('Accuracy Values')

# Combine Both Lines Into One Plot
lines = loss_line + accuracy_line
labels = [l.get_label() for l in lines]
ax1.legend(lines, labels, loc='lower right')
```

```
Out[20]: <matplotlib.legend.Legend at 0x2719c61b880>
```

